

Nishant Bhargava

bharga57@purdue.edu | linkedin.com/in/nishant-bhargava97 | github.com/nishant7309 |
nishant7309.github.io/portfolio

EDUCATION

Purdue University

GPA: 3.47/4.0

B.S. in Computer Engineering; Minors in Statistics

Expected May 2028

- **Relevant Coursework:** Multivariate Calculus, Linear Algebra, Statistics, Python for Data Science

PUBLICATION

Bhargava, N., et al. "Probe-Rewrite-Evaluate: Quantifying Evaluation Awareness in Large Language Models."
NeurIPS 2025 Workshop on Evaluation for Language Models & Wordplay Workshop @ EMNLP 2025.

[arXiv: 2509.00591](#)

EXPERIENCE

Incoming Summer 2026 AI Engineering Intern

Jun 2026 – Aug 2026

Empire State Realty Trust (ESRT)

New York, NY

- Incoming intern on the Application Development team, focused on building **Agentic AI** virtual assistant features for leasing and operational workflows.
- Will support integration of **Microsoft Copilot**, **ChatGPT**, and related AI tools into internal systems, with emphasis on testing, debugging, and documentation of AI workflows.

Undergraduate Research Assistant

Sep 2025 – Present

CYNICS Lab @ Purdue

West Lafayette, IN

- Developed a **Transformer**-based linguistic steganography pipeline that perturbs next-token distributions with **<2.0 PPL impact** while embedding encrypted payloads.
- Optimized multi-GPU **LLM** inference on **Modal** with **vLLM**, improving generation throughput by **~380%** in tokens/sec and accelerating internal experimentation cycles.
- **Fine-tuned** domain **LLMs** for **finance** and **pharma**, evaluating perplexity, task accuracy, and leakage risk while building a password-gated stego mechanism to securely serve authorized model variants.
- Investigated steganographic techniques on **Transformer-based semantic search**, demonstrating persistence of hidden text through vector embedding transformations.

AI Safety & Interpretability Researcher

May 2025 – Nov 2025

Algoverse

Remote

- Developed *Probe-Rewrite-Evaluate* methodology to quantify evaluation awareness in **LLMs**, revealing how models alter behavior between test-like and deployment contexts (NeurIPS '25 & EMNLP '25 workshop paper).
- Introduced **Awareness Elasticity** metric to quantify **LLM** sensitivity to evaluation interventions, providing the first systematic framework for measuring benchmark-deployment performance gaps.

Computer Vision Engineer

Jan 2025 – May 2025

RoBoat: Autonomous Maritime Maneuvers

Indianapolis, IN

- Led the AI and perception team for Purdue Indy's autonomous maritime robotics project in the AIMM ICC competition.
- Implemented deep learning-based object detection models (**YOLOv5**) to identify buoys, docks, and obstacles, achieving **93% detection accuracy** in simulation and 88%+ in real-world lake tests.
- Reduced false positives by **34%** through dataset curation, augmentation, and transfer learning on maritime-specific image data.

TECHNICAL SKILLS

Languages: Python, R, C, SQL, Bash, MATLAB

ML/LLM: PyTorch, scikit-learn, Hugging Face Transformers, transformer-lens, PEFT/LoRA, NumPy, pandas, OpenCV, Weights & Biases

Systems/Tools: Linux, Docker, Modal, AWS (SageMaker, S3, EC2), Git, vLLM, TensorRT